

R Package deSolve, Writing Code in Compiled Language

Karline Soetaert
Centre for
Estuarine and Marine Ecology
Netherlands Institute of Ecology
The Netherlands

Thomas Petzoldt
Technische Universität
Dresden
Germany

R. Woodrow Setzer
National Center for
Computational Toxicology
US Environmental Protection Agency

Abstract

This document describes how to use the **deSolve** package (Soetaert, Petzoldt, and Setzer 2009) to solve models that are written in FORTRAN or C.

Keywords: differential equation solvers, compiled code, performance, FORTRAN, C.

1. Introduction

deSolve (Soetaert *et al.* 2009), the successor of R package **odesolve** (Setzer 2001) is a package to solve ordinary differential equations (ODE), differential algebraic equations (DAE) and partial differential equations (PDE).

One of the prominent features of **deSolve** is that it allows specifying the differential equations either as:

- pure R code (R Development Core Team 2008),
- functions defined in lower-level languages such as FORTRAN, C, or C++, which are compiled into a dynamically linked library (DLL) and loaded into R.

In what follows, these implementations will be referred to as **R models** and **DLL models** respectively.

Whereas **R models** are easy to implement, they allow simple interactive development, produce highly readable code and access to R's high-level procedures, **DLL models** have the benefit of increased simulation speed. Depending on the problem, there may be a gain of up to several orders of magnitude computing time when using compiled code.

Here are some rules of thumb when it is worthwhile or not to switch to **DLL models**.

- As long as one makes use only of R's high-level commands, the time gain will be modest. This was demonstrated in Soetaert, Petzoldt, and Setzer (subm), where a formulation of two interacting populations dispersing on a 1-dimensional or a 2-dimensional grid led to a time gain of a factor two only when using **DLL models**.

- Generally, the more statements in the model, the higher will be the gain of using compiled code. Thus, in the same paper (Soetaert et al. subm.), a very simple, 0-D, Lotka-Volterra type of model describing only 2 state variables was solved 50 times faster when using compiled code.
- As even **R models** are quite performant, the time gain induced by compiled code will often not be discernible when the model is only solved once (who can grasp the difference between a run taking 0.001 or 0.05 seconds to finish). However, if the model is to be applied multiple times, e.g. because the model is to be fitted to data, or its sensitivity is to be tested, then it may be worthwhile to implement the model in a compiled language.

Starting from deSolve version 1.4, it is now also possible to use *forcing functions* in compiled code. These forcing functions are automatically updated by the integrators. See last chapter.

2. A simple ODE example

Assume the following simple ODE (which is from the LSODA source code):

$$\begin{aligned}\frac{dy_1}{dt} &= -k_1 \cdot y_1 + k_2 \cdot y_2 \cdot y_3 \\ \frac{dy_2}{dt} &= k_1 \cdot y_1 - k_2 \cdot y_2 \cdot y_3 - k_3 \cdot y_2 \cdot y_2 \\ \frac{dy_3}{dt} &= k_3 \cdot y_2 \cdot y_2\end{aligned}$$

where y_1 , y_2 and y_3 are state variables, and k_1 , k_2 and k_3 are parameters.

We first implement and run this model in pure R, then show how to do this in C and in FORTRAN.

2.1. ODE model implementation in R

An ODE model implemented in **pure R** should be defined as:

```
yprime = func(t, y, parms,...)
```

where **t** is the current time point in the integration, **y** is the current estimate of the variables in the ODE system, and **parms** is a vector or list containing the parameter values. ... (optional) are any other arguments passed to the function. The return value of **func** should be a list, whose first element is a vector containing the derivatives of **y** with respect to time, and whose next elements contain output variables that are required at each point in time.

The R implementation of the simple ODE is given below:

```
R> model <- function(t, Y, parameters) {
+   with(as.list(parameters), {
```

```

+      dy1 = -k1 * Y[1] + k2 * Y[2] * Y[3]
+      dy3 = k3 * Y[2] * Y[2]
+      dy2 = -dy1 - dy3
+      list(c(dy1, dy2, dy3))
+    })
+  }

```

The Jacobian ($\frac{\partial y'}{\partial y}$) associated to the above example is:

```

R> jac <- function(t, Y, parameters) {
+   with(as.list(parameters), {
+     PD[1, 1] = -k1
+     PD[1, 2] = k2 * Y[3]
+     PD[1, 3] = k2 * Y[2]
+     PD[2, 1] = k1
+     PD[2, 3] = -PD[1, 3]
+     PD[3, 2] = k3 * Y[2]
+     PD[2, 2] = -PD[1, 2] - PD[3, 2]
+     return(PD)
+   })
+ }

```

This model can then be run as follows:

```

R> parms <- c(k1 = 0.04, k2 = 10000, k3 = 3e+07)
R> Y <- c(1, 0, 0)
R> times <- c(0, 0.4 * 10^(0:11))
R> PD <- matrix(nrow = 3, ncol = 3, data = 0)
R> out <- ode(Y, times, model, parms = parms, jacfunc = jac)

```

2.2. ODE model implementation in C

In order to create compiled models (.DLL = dynamic link libraries on Windows or .so = shared objects on other systems) you must have a recent version of the GNU compiler suite installed, which is quite standard for Linux. Windows users find all the required tools on <http://www.murdoch-sutherland.com/Rtools/>. Getting DLLs produced by other compilers to communicate with R is much more complicated and therefore not recommended. More details can be found on <http://cran.r-project.org/doc/manuals/R-admin.html>.

The call to the derivative and Jacobian function is more complex for compiled code compared to R-code, because it has to comply with the interface needed by the integrator source codes. Below is an implementation of this model in C:

```

/* file mymod.c */
#include <R.h>
static double parms[3];

```

```

#define k1 parms[0]
#define k2 parms[1]
#define k3 parms[2]

/* initializer */
void initmod(void (* odeparms)(int *, double *))
{
    int N=3;
    odeparms(&N, parms);
}

/* Derivatives and 1 output variable */
void derivs (int *neq, double *t, double *y, double *ydot,
             double *yout, int *ip)
{
    if (ip[0] <1) error("nout should be at least 1");
    ydot[0] = -k1*y[0] + k2*y[1]*y[2];
    ydot[2] = k3 * y[1]*y[1];
    ydot[1] = -ydot[0]-ydot[2];

    yout[0] = y[0]+y[1]+y[2];
}

/* The Jacobian matrix */
void jac(int *neq, double *t, double *y, int *ml, int *mu,
         double *pd, int *nrowpd, double *yout, int *ip)
{
    pd[0]          = -k1;
    pd[1]          = k1;
    pd[2]          = 0.0;
    pd[(*nrowpd)]  = k2*y[2];
    pd[(*nrowpd) + 1] = -k2*y[2] - 2*k3*y[1];
    pd[(*nrowpd) + 2] = 2*k3*y[1];
    pd[(*nrowpd)*2] = k2*y[1];
    pd[2*(*nrowpd) + 1] = -k2 * y[1];
    pd[2*(*nrowpd) + 2] = 0.0;
}
/* END file mymod.c */

```

The implementation in C consists of three parts:

1. After defining the parameters in global C-variables, through the use of **#define** statements, a function called **initmod** initialises the parameter values, passed from the R-code.

This function has as its sole argument a pointer to C-function **odeparms** that fills a double array with double precision values, to copy the parameter values into the global variable.

2. Function `derivs` then calculates the values of the derivatives. The derivative function is defined as:

```
void derivs (int *neq, double *t, double *y, double *ydot,
            double *yout, int *ip)
```

where `*neq` is the number of equations, `*t` is the value of the independent variable, `y` points to a double precision array of length `*neq` that contains the current value of the state variables, and `ydot` points to an array that will contain the calculated derivatives.

`yout` points to a double precision vector whose first `nout` values are other output variables (different from the state variables `y`), and the next values are double precision values as passed by parameter `rpar` when calling the integrator. The key to the elements of `yout` is set in `*ip`

`*ip` points to an integer vector whose length is at least 3; the first element (`ip[0]`) contains the number of output values (which should be equal or larger than `nout`), its second element contains the length of `*yout`, and the third element contains the length of `*ip`; next are integer values, as passed by parameter `ipar` when calling the integrator.¹

Note that, in function `derivs`, we start by checking whether enough room is allocated for the output variables (`if (ip[0] < 1)`), else an error is passed to R and the integration is stopped.

3. In C, the call to the function that generates the Jacobian is as:

```
void jac(int *neq, double *t, double *y, int *ml,
        int *mu, double *pd, int *nrowpd, double *yout, int *ip)
```

where `*ml` and `*mu` are the number of non-zero bands below and above the diagonal of the Jacobian respectively. These integers are only relevant if the option of a banded Jacobian is selected. `*nrow` contains the number of rows of the Jacobian. Only for full Jacobian matrices, this is equal to `*neq`. In case the Jacobian is banded, `*nrowpd` will be equal to `*mu+*ml+1`.²

2.3. ODE model implementation in FORTRAN

Models may also be defined in FORTRAN.

¹Readers familiar with the source code of the **ODEPACK** solvers may be surprised to find the double precision vector `yout` and the integer vector `ip` at the end. Indeed none of the **ODEPACK** functions allow this, although it is standard in the `ode` and `daspk` codes. To make all integrators compatible, we have altered the **ODEPACK** FORTRAN codes to consistently pass these vectors.

²Readers familiar with the implementation of the FORTRAN code `DVODE` may notice that this is not the format in which `DVODE` requires the specification of the Jacobian; this code needs an extra `mu` empty rows. As we have taken the philosophy to make the model specification independent of the integrator that will be used, the facility to use `ode` with a Jacobian specified in compiled code has been toggled off. Use the related code `lsode` instead.

```

c file mymodf.f
  subroutine initmod(odeparms)
    external odeparms
    double precision parms(3)
    common /myparms/parms

    call odeparms(3, parms)
    return
  end

  subroutine derivs (neq, t, y, ydot, yout, ip)
    double precision t, y, ydot, k1, k2, k3
    integer neq, ip(*)
    dimension y(3), ydot(3), yout(*)
    common /myparms/k1,k2,k3

    if(ip(1) < 1) call rexit("nout should be at least 1")

    ydot(1) = -k1*y(1) + k2*y(2)*y(3)
    ydot(3) = k3*y(2)*y(2)
    ydot(2) = -ydot(1) - ydot(3)

    yout(1) = y(1) + y(2) + y(3)
    return
  end

  subroutine jac (neq, t, y, ml, mu, pd, nrowpd, yout, ip)
    integer neq, ml, mu, nrowpd, ip
    double precision y(*), pd(nrowpd,*), yout(*), t, k1, k2, k3
    common /myparms/k1, k2, k3

    pd(1,1) = -k1
    pd(2,1) = k1
    pd(3,1) = 0.0
    pd(1,2) = k2*y(3)
    pd(2,2) = -k2*y(3) - 2*k3*y(2)
    pd(3,2) = 2*k3*y(2)
    pd(1,3) = k2*y(2)
    pd(2,3) = -k2*y(2)
    pd(3,3) = 0.0
    return
  end
c end of file mymodf.f

```

In FORTRAN, parameters may be stored in a common block (here called `myparms`). During the initialisation, this common block is defined to consist of a 3-valued vector (unnamed), but in the subroutines `derivs` and `jac`, the parameters are given a name (`k1,...`).

2.4. Running ODE models implemented in compiled code

To run the models described above, the code in `mymod.f` and `mymod.c` must first be compiled³. This can simply be done in R itself, using the `system` command:

```
R> system("R CMD SHLIB mymod.f")
```

for the FORTRAN code or

```
R> system("R CMD SHLIB mymod.c")
```

for the C code.

This will create file `mymod.dll`. After loading the DLL, the model can be run:

```
R> dyn.load("mymod.dll")
R> parms <- c(k1 = 0.04, k2 = 10000, k3 = 3e+07)
R> Y <- c(y1 = 1, y2 = 0, y3 = 0)
R> times <- c(0, 0.4 * 10^(0:11))
R> out <- ode(Y, times, func = "derivs", parms = parms,
+   jacfunc = "jac", dllname = "mymod", initfunc = "initmod",
+   nout = 1, outnames = "Sum")
```

The integration routine (here `ode`) recognizes that the model is specified as a DLL due to the fact that arguments `func` and `jacfunc` are not regular R-functions but character strings. Thus, the integrator will check whether the function is loaded in the DLL with name "mymod". Note that "mymod", as specified by `dllname` gives the name of the shared library *without extension*. This DLL should contain all the compiled function or subroutine definitions referred to in `func`, `jacfunc` and `initfunc`.

Also, if `func` is specified in compiled code, then `jacfunc` and `initfunc` (if present) should also be specified in a compiled language. It is not allowed to mix R-functions and compiled functions

Note also that, when invoking the integrator, we have to specify the number of ordinary output variables, `nout`. This is because the integration routine has to allocate memory to pass these output variables back to R. There is no way to check for the number of output variables in a DLL automatically. If in the calling of the integration routine the number of output variables is too low, then R may freeze and need to be terminated! Therefore it is advised that one checks in the code whether `nout` has been specified correctly. In the FORTRAN example above, the statement `if (ip(1) < 1) call rexit("nout should be at least 1")` does this. Note that it is not an error (just a waste of memory) to set `nout` to a too large value.

Finally, in order to label the output matrix, the name of the ordinary output variable has to be passed explicitly (`outnames`). The names of the state variables are known through their initial condition (`y`).

³This requires a correctly installed GNU compiler, see above.

3. Passing parameters and data in compiled code

All of the solvers in **deSolve** take an argument **parms** which may be an arbitrary R object. In models defined in R code, this argument is passed unprocessed to the various functions that make up the model. It is possible, as well, to pass such objects to models defined in native code.

The problem is that data passed to, say, **ode** in the argument **parms** is not visible by default to the routines that define the model. This is handled by a user-written initialization function, for example **initmod** in the C and FORTRAN examples from sections 2.2 and 2.3. In C, the initialization routine is declared:

```
void initmod(void (* odeparms)(int *, double *));
```

That is, **initmod** has a single argument, a pointer to a function that has as arguments a pointer to an **int** and a pointer to a **double**. In FORTRAN, the initialization routine has a single argument, a subroutine declared to be external. The name of the initialization function is passed as an argument to the **deSolve** solver functions.

In C, two approaches are available for making the values passed in **parms** visible to the model routines, while only the simpler approach is available in FORTRAN. The simpler requires that **parms** be a numeric vector. In C, the function passed from **deSolve** to the initialization function (called **odeparms** in the example) copies the values from the parameter vector to a static array declared globally in the file where the model is defined. In FORTRAN, the values are copied into a **COMMON** block.

It is possible to pass more complicated structures to C functions. Here is an example, an initializer called **deltamethrin** from a model describing the pharmacokinetics of that pesticide:

```
#include <R.h>
#include <Rinternals.h>
#include <R_ext/Rdynload.h>
#include "deltamethrin.h"

/* initializer */
void deltamethrin(void(* odeparms)(int *, double *))
{
    int Nparms;
    DL_FUNC get_deSolve_gparms;
    SEXP gparms;
    get_deSolve_gparms = R_GetCCallable("deSolve","get_deSolve_gparms");
    gparms = get_deSolve_gparms();
    Nparms = LENGTH(gparms);
    if (Nparms != N_PARMS) {
        PROBLEM "Confusion over the length of parms"
        ERROR;
    } else {
        _RDy_deltamethrin_parms = REAL(gparms);
    }
}
```


In `deltamethrin.h`, the variable `_RDy_deltamethrin_parms` and macro `N_PARMS` are declared:

```
#define N_PARMS 63
static double *_RDy_deltamethrin_parms;
```

The critical element of this method is the function `R_GetCCallable` which returns a function (called `get_deSolve_gparms` in this implementation) that returns the `parms` argument as a `SEXP` data type. In this example, `parms` was just a real vector, but in principle, this method can handle arbitrarily complex objects. For more detail on handling R objects in native code, see R Development Core Team (2008).

4. deSolve integrators that support DLL models

In the most recent version of **deSolve** all integration routines can solve **DLL models**. They are:

- all solvers of the `lsode` family: `lsoda`, `lsode`, `lsodar`, `lsodes`
- `vode`
- `daspk`
- the C-implemented Runge-Kutta integration routines (including the Euler method)

For some of these solvers the interface is slightly different (e.g. `daspk`), while in others (`lsodar`, `lsodes`) different functions can be defined. How this is implemented in a compiled language is discussed next.

4.1. DAE models, integrator `daspk`

`daspk` is the only integrator in the package that solves DAE models. DAEs are specified in implicit form:

$$0 = F(y', y, x, t)$$

i.e. the DAE function (passed via argument `res`) specifies the “residuals” rather than the derivatives (as for ODEs).

Consequently the DAE function specification in compiled language are also different. For code written in C, the calling sequence for `res` must be:

```
void myres(double *t, double *y, double *ydot, double *cj,
           double *delta, int *ires, double *yout, int *ip)
```

where `*t` is the value of the independent variable, `y` points to a double precision array that contains the current value of the state variables, `ydot` points to an array that will contain the derivatives, `delta` points to an array that will contain the calculated residuals. `cj` points to a scalar, which is normally proportional to the inverse of the stepsize, while `ires` points

to an integer (not used). `yout` points to any other output variables (different from the state variables `y`), followed by the double precision values as passed via argument `rpar`; finally `*ip` is an integer vector containing at least 3 elements, its first value (`*ip[0]`) equals the number of output variables, calculated in the function (and which should be equal to `nout`), its second element equals the total length of `*yout`, its third element equals the total length of `*ip`, and finally come the integer values as passed via argument `ipar`.

For code written in FORTRAN, the calling sequence for `res` must be as in the following example:

```

subroutine myresf(t, y, ydot, cj, delta, ires, out, ip)
  integer :: ires, ip(*)
  integer, parameter :: neq = 3
  double precision :: t, y(neq), ydot(neq), delta(neq), out(*)
  double precision :: K, ka, r, prod, ra, rb
  common /myparms/K,ka,r,prod

  if(ip(1) < 1) call rexit("nout should be at least 1")
  ra = ka* y(3)
  rb = ka/K *y(1) * y(2)

!! residuals of rates of changes
  delta(3) = -ydot(3) - ra + rb + prod
  delta(1) = -ydot(1) + ra - rb
  delta(2) = -ydot(2) + ra - rb - r*y(2)
  out(1) = y(1) + y(2) + y(3)
  return
end

```

Similarly as for the ODE model discussed above, the parameters are kept in a common block which is initialised by an initialiser subroutine:

```

subroutine initpar(daspkparms)

  external daspkparms
  integer, parameter :: N = 4
  double precision parms(N)
  common /myparms/parms
  call daspkparms(N, parms)
  return
end

```

See the ODE example for how to initialise parameter values in C.

Similarly, the function that specifies the Jacobian in a DAE differs from the Jacobian when the model is an ODE. The DAE Jacobian is set with argument `jacres` rather than `jacfunc` when an ODE.

For code written in FORTRAN, the `jacres` must be as:

```

subroutine resjacfor (t, y, dy, pd, cj, out, ipar)

  integer, parameter :: neq = 3
  integer :: ipar(*)
  double precision :: K, ka, r, prod
  double precision :: pd(neq,neq),y(neq),dy(neq),out(*)
  common /myparms/K,ka,r,prod

!res1 = -dD - ka*D + ka/K *A*B + prod
  PD(1,1) = ka/K *y(2)
  PD(1,2) = ka/K *y(1)
  PD(1,3) = -ka -cj
!res2 = -dA + ka*D - ka/K *A*B
  PD(2,1) = -ka/K *y(2) -cj
  PD(2,2) = -ka/K *y(2)
  PD(2,3) = ka
!res3 = -dB + ka*D - ka/K *A*B - r*B
  PD(3,1) = -ka/K *y(2)
  PD(3,2) = -ka/K *y(2) -r -cj
  PD(3,3) = ka
  return
end

```

4.2. The root function from integrator lsodar

lsodar is an extended version of integrator **lsoda** that includes a root finding function. This function is specified via argument **rootfunc**.

Here is how to program such a function in a lower-level language. For code written in C, the calling sequence for **rootfunc** must be:

```

void myroot(int *neq, double *t, double *y, int *ng, double *gout,
            double *out, int *ip )

```

where ***neq** and ***ng** are the number of state variables and root functions respectively, ***t** is the value of the independent variable, **y** points to a double precision array that contains the current value of the state variables, and **gout** points to an array that will contain the values of the constraint function whose root is sought. ***out** and ***ip** are a double precision and integer vector respectively, as described in the ODE example above.

For code written in **FORTRAN**, the calling sequence for **rootfunc** must be as in following example:

```

subroutine myroot(neq, t, y, ng, gout, out, ip)
  integer :: neq, ng, ip(*)
  double precision :: t, y(neq), gout(ng), out(*)

  gout(1) = y(1) - 1.e-4

```

```

gout(2) = y(3) - 1e-2

return
end

```

4.3. jacvec, the Jacobian vector for integrator lsodes

Finally, in integration function `lsodes`, not the Jacobian matrix is specified, but a vector, one for each column of the Jacobian. This function is specified via argument `jacvec`.

In FORTRAN, the calling sequence for `jacvec` is:

```

SUBROUTINE JAC (NEQ, T, Y, J, IAN, JAN, PDJ, OUT, IP)
DOUBLE PRECISION T, Y(*), IAN(*), JAN(*), PDJ(*), OUT(*)
INTEGER NEQ, J, IP(*)

```

5. Testing functions written in compiled code

Two utilities have been included to test the function implementation in compiled code:

- **DLLfunc** to test the implementation of the derivative function as used in ODEs. This function returns the derivative $\frac{dy}{dt}$ and the output variables.
- **DLLres** to test the implementation of the residual function as used in DAEs. This function returns the residual function $\frac{dy}{dt} - f(y, t)$ and the output variables.

These functions serve no other purpose than to test whether the compiled code returns what it should.

5.1. DLLfunc

We test whether the `ccl4` model, which is part of **deSolve** package, returns the proper rates of changes. (Note: see `example(ccl4model)` for a more comprehensive implementation)

```

R> Params <- c(0.182, 4, 4, 0.08, 0.04, 0.74, 0.05, 0.15,
+ 0.32, 16.17, 281.48, 13.3, 16.17, 5.487, 153.8, 0.04321671,
+ 0.4027255, 1000, 0.02, 1, 3.8)
R> yini <- c(AI = 21, AAM = 0, AT = 0, AF = 0, AL = 0, CLT = 0,
+ AM = 0)
R> DLLfunc(y = yini, dllname = "deSolve", func = "derivsccl4",
+ initfunc = "initccl4", parms = Params, times = 1,
+ nout = 3, outnames = c("DOSE", "MASS", "CP"))

```

\$dy

AI	AAM	AT	AF	AL
-20.0582048	6.2842256	9.4263383	0.9819102	2.9457307

```

          CLT          AM
0.0000000  0.0000000

$var
      DOSE      MASS      CP
1.758626  0.000000 922.727067

```

5.2. DLLres

The deSolve package contains a FORTRAN implementation of the chemical model described above (section 4.1), where the production rate is included as a forcing function (see next section).

Here we use DLLres to test it:

```

R> pars <- c(K = 1, ka = 1e+06, r = 1)
R> y <- c(A = 2, B = 3, D = 2 * 3/pars["K"])
R> dy <- c(dA = 0, dB = 0, dD = 0)
R> prod <- matrix(nc = 2, data = c(seq(0, 100, by = 10),
+   seq(0.1, 0.5, len = 11)))
R> DLLres(y = y, dy = dy, times = 5, res = "chemres", dllname = "deSolve",
+   initfunc = "initparms", initforc = "initforcs", parms = pars,
+   forcings = prod, nout = 2, outnames = c("CONC", "Prod"))

$delta
      A      B  D.K
0.00 -3.00  0.12

$var
  CONC  Prod
11.00  0.12

```

6. Using forcing functions

Forcing functions in DLLs are implemented in a similar way as parameters. This means:

- They are initialised by means of an initialiser function. Its name should be passed to the solver via argument `initforc`.

Similar as the parameter initialiser function, the function denoted by `initforc` has as its sole argument a pointer to the vector that contains the forcing function values in the compiled code. In case of C code, this will be a global vector; in case of FORTRAN, this will be a vector in a common block.

The solver puts a pointer to this vector and updates the forcing functions in this memory area at each time step. Hence, within the compiled code, forcing functions can be assessed as if they are parameters (although, in contrast to the latter, their values will

generally change). No need to update the values for the current time step; this has been done before entering the `derivs` function.

- The forcing function data series are passed to the integrator, via argument `forcings`; if there is only one forcing function data set, then a 2-columned matrix (time, value) will do; else the data should be passed as a list, containing (time, value) matrices with the individual forcing function data sets. Note that the data sets in this list should be *in the same ordering* as the declaration of the forcings in the compiled code.

A number of options allow to finetune certain settings. They are in a list called `fcontrol` which can be supplied as argument when calling the solvers. The options are similar to the arguments from R function `approx`, however the default settings are often different.

The following options can be specified:

- `method` specifies the interpolation method to be used. Choices are "linear" or "constant", the default is "linear", which means linear interpolation (same as `approx`)
- `rule`, an integer describing how interpolation is to take place *outside* the interval $[\min(\text{times}), \max(\text{times})]$. If `rule` is 1 then an error will be triggered and the calculation will stop if interpolation is necessary. If it is 2, the default, the value at the closest data extreme is used, a warning will be printed if `verbose` is TRUE.

Note that the default differs from the `approx` default.

- `f`, for `method="constant"` is a number between 0 and 1 inclusive, indicating a compromise between left- and right-continuous step functions. If `y0` and `y1` are the values to the left and right of the point then the value is $y0*(1-f)+y1*f$ so that `f=0` is right-continuous and `f=1` is left-continuous. The default is to have `f=0`. For some data sets it may be more realistic to set `f=0.5`.
- `ties`, the handling of tied `times` values. Either a function with a single vector argument returning a single number result or the string "ordered".

Note that the default is "ordered", hence the existence of ties will NOT be investigated; in practice this means that, if ties exist, the first value will be used; if the dataset is not ordered, then nonsense will be produced.

Alternative values for `ties` are `mean`, `min` etc... which will average, or take the minimal value if multiple values exist at one time level.

The default settings of `fcontrol` are:

```
fcontrol=list(method="linear", rule = 2, f = 0, ties = "ordered")
```

Note that only ONE specification is allowed, even if there is more than one forcing function data set. (may/should change in the future).

6.1. A simple FORTRAN example

We implement the example from chapter 3 of the book ([Soetaert and Herman 2009](#)) in FORTRAN.

This model describes the oxygen consumption of a (marine) sediment in response to deposition of organic matter (the forcing function). One state variable, the organic matter content in the sediment is modeled; it changes as a function of the deposition **Flux** (forcing) and organic matter decay (first-order decay rate **k**).

$$\frac{dC}{dt} = Flux_t - k \cdot C$$

with initial condition $C(t = 0) = C_0$; the latter is estimated as the mean of the flux divided by the decay rate.

The FORTRAN code looks like this:

```
c Initialiser for parameter common block
  subroutine scocpar(odeparms)

    external odeparms
    integer N
    double precision parms(2)
    common /myparms/parms

    N = 1
    call odeparms(N, parms)
    return
  end

c Initialiser for forcing common block
  subroutine scocforc(odeforcs)

    external odeforcs
    integer N
    double precision forcs(1)
    common /myforcs/forcs

    N = 1
    call odeforcs(N, forcs)
    return
  end

c Rate of change and output variables

  subroutine scocder (neq, t, y, ydot,out,IP)

    integer          neq, IP(*)
    double precision t, y(neq), ydot(neq), out(*), k, depo
    common /myparms/k
    common /myforcs/depo
```

```

if(IP(1) < 2) call rexit("nout should be at least 2")

ydot(1) = -k*y(1) + depo

out(1)= k*y(1)
out(2)= depo

return
end

```

Here the subroutine `scocpar` is business as usual; it initialises the parameter common block (there is only one parameter). Subroutine `odeforcs` does the same for the forcing function, which is also positioned in a common block, called `myforcs`. This common block is made available in the derivative subroutine (here called `scocder`), where the forcing function is named `depo`.

At each time step, the integrator updates the value of this forcing function to the correct time point. In this way, the forcing functions can be used as if they are (time-varying) parameters. All that's left to do is to pass the forcing function data set and the name of the forcing function initialiser routine. This is how to do it in R.

First the data are inputted:

```

R> Flux <- matrix(ncol = 2, byrow = TRUE, data = c(1, 0.654,
+ 11, 0.167, 21, 0.06, 41, 0.07, 73, 0.277, 83, 0.186,
+ 93, 0.14, 103, 0.255, 113, 0.231, 123, 0.309, 133,
+ 1.127, 143, 1.923, 153, 1.091, 163, 1.001, 173, 1.691,
+ 183, 1.404, 194, 1.226, 204, 0.767, 214, 0.893, 224,
+ 0.737, 234, 0.772, 244, 0.726, 254, 0.624, 264, 0.439,
+ 274, 0.168, 284, 0.28, 294, 0.202, 304, 0.193, 315,
+ 0.286, 325, 0.599, 335, 1.889, 345, 0.996, 355, 0.681,
+ 365, 1.135))
R> head(Flux)

```

```

      [,1] [,2]
[1,]    1 0.654
[2,]   11 0.167
[3,]   21 0.060
[4,]   41 0.070
[5,]   73 0.277
[6,]   83 0.186

```

and the parameter given a value (there is only one)

```

R> parms <- 0.01

```

The initial condition `Yini` is estimated as the annual mean of the Flux and divided by the decay rate (parameter).


```
R> meanDepo <- mean(approx(Flux[, 1], Flux[, 2], xout = seq(1,
+ 365, by = 1))$y)
R> Yini <- c(y = meanDepo/parms)
```

After defining the output times, the model is run, using integration routine `ode`.

The *name* of the derivate function "`scocder`", of the dll "`deSolve`"⁴ and of the initialiser function "`scocpar`" are passed, as in previous examples.

In addition, the forcing function data set is also passed (`forcings=Flux`) as is the name of the forcing initialisation function (`initforc="scocforc"`).

```
R> times <- 1:365
R> out <- as.data.frame(ode(y = Yini, times, func = "scocder",
+ parms = parms, dllname = "deSolve", initforc = "scocforc",
+ forcings = Flux, initfunc = "scocpar", nout = 2,
+ outnames = c("Mineralisation", "Depo")))
R> head(out)
```

	time	y	Mineralisation	Depo
1	1	63.00301	0.6300301	0.6540
2	2	63.00262	0.6300262	0.6053
3	3	62.95377	0.6295377	0.5566
4	4	62.85694	0.6285694	0.5079
5	5	62.71259	0.6271259	0.4592
6	6	62.52124	0.6252124	0.4105

Now, the way the forcing functions are interpolated are changed: Rather than linear interpolation, constant (block, step) interpolation is used.

```
R> fcontrol <- list(method = "constant")
R> out2 <- as.data.frame(ode(y = Yini, times, func = "scocder",
+ parms = parms, dllname = "deSolve", initforc = "scocforc",
+ forcings = Flux, fcontrol = fcontrol, initfunc = "scocpar",
+ nout = 2, outnames = c("Mineralisation", "Depo")))
```

Finally, the results are plotted:

```
R> par(mfrow = c(1, 2))
R> plot(out$time, out$Depo, type = "l", col = "red", xlab = "days",
+ ylab = "mmol C/m2/ d", main = "method='linear'")
R> lines(out$time, out$Mineralisation, lwd = 2, col = "blue")
R> legend("topleft", lwd = 1:2, col = c("red", "blue"),
+ c("Flux", "Mineralisation"))
R> plot(out2$time, out2$Depo, type = "l", col = "red", xlab = "days",
+ ylab = "mmol C/m2/ d", main = "method='constant'")
R> lines(out2$time, out2$Mineralisation, lwd = 2, col = "blue")
```




Figure 1: Solution of the SCOC model, implemented in compiled code, and including a forcing function - see text for R-code

6.2. An example in C

Consider the following R-code which implements a resource-producer-consumer Lotka-Volterra type of model in R (it is a modified version of the example of function `ode`):

```
R> SPCmod <- function(t, x, parms, input) {
+   with(as.list(c(parms, x)), {
+     import <- input(t)
+     dS <- import - b * S * P + g * C
+     dP <- c * S * P - d * C * P
+     dC <- e * P * C - f * C
+     res <- c(dS, dP, dC)
+     list(res, signal = import)
+   })
+ }
R> parms <- c(b = 0.1, c = 0.1, d = 0.1, e = 0.1, f = 0.1,
+   g = 0)
R> times <- seq(0, 100, by = 0.1)
R> signal <- as.data.frame(list(times = times, import = rep(0,
+   length(times))))
R> signal$import <- ifelse((trunc(signal$times)%2 == 0),
+   0, 1)
R> sigimp <- approxfun(signal$times, signal$import, rule = 2)
R> xstart <- c(S = 1, P = 1, C = 1)
R> print(system.time(out <- ode(y = xstart, times = times,
+   func = SPCmod, parms, input = sigimp)))
```

user	system	elapsed
1.10	0.04	1.12

⁴this example is made part of the `deSolve` package, hence the name of the dll is "deSolve"

All output is printed at once:

```
R> plot(out, type = "l")
```

The C-code, in file `Forcing_lv.c`, can be found in the packages `/inst/dynload` subdirectory. It can be compiled, from within R by

```
system("R CMD SHLIB Forcing_lv.c")
```

After defining the parameter and forcing vectors, and giving them comprehensible names, the parameter and forcing initialiser functions are defined (`parmsc` and `forcc` respectively). Next is the derivative function, `derivsc`.

```
#include <R.h>

static double parms[6];
static double forc[1];

/* A trick to keep up with the parameters and forcings */
#define b parms[0]
#define c parms[1]
#define d parms[2]
#define e parms[3]
#define f parms[4]
#define g parms[5]

#define import forc[0]

/* initializers: */
void odec(void (* odeparms)(int *, double *))
{
    int N=6;
    odeparms(&N, parms);
}

void forcc(void (* odeforcs)(int *, double *))
{
    int N=1;
    odeforcs(&N, forc);
}

/* derivative function */
void derivsc(int *neq, double *t, double *y, double *ydot,
             double *yout, int*ip)
{
    if (ip[0] <2) error("nout should be at least 2");
```




Figure 2: Solution of the lotka-volterra resource (S)-producer (P) - consumer (C) model with time-variable input (signal) - see text for R-code

```

ydot[0] = import - b*y[0]*y[1] + g*y[2]      ;
ydot[1] = c*y[0]*y[1] - d*y[2]*y[1]         ;
ydot[2] = e*y[1]*y[2] - f*y[2]              ;

yout[0] = y[0]+y[1]+y[2];
yout[1] = import;
}

```

After defining the forcing function time series, which is to be interpolated by the integration routine, and loading the DLL, the model is run:

```

Sigimp <- approx(signal$times, signal$import, xout=ftime,rule = 2)$y
forcings <- cbind(ftime,Sigimp)

dyn.load("Forcing_lv.dll")
out <- ode(y=xstart, times, func = "derivsc",
  parms = parms, dllname = "Forcing_lv",initforc="forcc",
  forcings=forcings, initfunc = "parmsc", nout = 2,
  outnames = c("Sum","signal"),method=rkMethod("rk34f"))
dyn.unload("Forcing_lv.dll")

```

This code executes about 30 times faster than the R-code.

With a longer simulation time, the difference becomes more pronounced, e.g. with times till 800 days, the DLL code executes 200 times faster⁵.

7. Final remark

Detailed information about communication between C, FORTRAN and R can be found in [R Development Core Team \(2009\)](#).

Notwithstanding the speed gain when using compiled code, one should not carelessly decide to always resort to this type of modelling.

Because the code needs to be formally compiled and linked to R much of the elegance when using pure R models is lost. Moreover, mistakes are easily made and paid harder in compiled code: often a programming error will terminate R. In addition, these errors may not be simple to trace.

⁵this is due to the sequential update of the forcing functions by the solvers, compared to the bisectioning approach used by `approxfun`

References

- R Development Core Team (2008). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-07-0, URL <http://www.R-project.org>.
- R Development Core Team (2009). *Writing R Extensions*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-11-9, URL <http://www.R-project.org>.
- Setzer RW (2001). *The **odesolve** Package: Solvers for Ordinary Differential Equations*. R package version 0.1-1.
- Soetaert K, Herman PMJ (2009). *A Practical Guide to Ecological Modelling. Using R as a Simulation Platform*. Springer. ISBN 978-1-4020-8623-6.
- Soetaert K, Petzoldt T, Setzer R (subm). “Solving Differential Equations in R: Package **deSolve**.” *Journal of Statistical Software*.
- Soetaert K, Petzoldt T, Setzer RW (2009). *deSolve: General solvers for initial value problems of ordinary differential equations (ODE), partial differential equations (PDE) and differential algebraic equations (DAE)*. R package version 1.4.

Affiliation:

Karline Soetaert
Centre for Estuarine and Marine Ecology (CEME)
Netherlands Institute of Ecology (NIOO)
4401 NT Yerseke, Netherlands
E-mail: k.soetaert@nioo.knaw.nl
URL: <http://www.nioo.knaw.nl/users/ksoetaert>

Thomas Petzoldt
Institut für Hydrobiologie
Technische Universität Dresden
01062 Dresden, Germany
E-mail: thomas.petzoldt@tu-dresden.de
URL: <http://tu-dresden.de/Members/thomas.petzoldt/>

R. Woodrow Setzer
National Center for Computational Toxicology
US Environmental Protection Agency
URL: <http://www.epa.gov/comptox>